

## **TASK 1**

1. In order to answer to this question, only the heroku/[router] and app/web.X log entries were used. According to heroku's documentation page, the inbound requests are received by a set of routers that forward the request to one of the application's web dynos. The web dynos are the only type of dynos that can receive HTTP requests. The router logs can be correlated with the web dynos logs using the request\_id field. Two different dataframes were created, one for the router logs and one for the web dynos logs. To count how many times the status=404 had appeared in each dataframe, it was ensured that the same request\_id is only counted once.

The not found url with the highest count is the :  
workabledemo.com/api/accounts/3

```
>>> not_found_urls_a
      0  counts
0  path=/api/accounts/3    558
>>> not_found_urls_r
      path      host  counts
0      /accounts  www.workabledemo.com    1
1  /api/accounts/3  workabledemo.com    4378
2  /backend/subscription/update_billing  sampleco.workabledemo.com    1
3      /petitions  www.workabledemo.com    1
4  /uas/request-password-reset?trk  www.workabledemo.com    1
5      /user_password_resets  www.workabledemo.com    1
>>>
```

2.

```
>>> avg_service=DF_router["service"].str.strip("ms").astype(int).mean()
>>> avg_service
3.83262980160405
>>>
```

3. The table “delayed jobs” is more frequently loaded.

```
sorted_DF=tables_DF_sum.sort_values(by=["sum"],ascending=False)
sorted_DF
```

	index	sum
	"delayed_jobs"	16741
	"accounts"	5417
	"jobs"	1426
	"members"	1216
	"collaborations"	911
	"candidates"	849
	"postings"	607
	"job_boards"	391
	"stages"	380
	"users"	301
	"subscriptions"	238
	"job_stats"	108
	"plans"	91
	"keywords"	89
	"slot_plans"	75
	"activities"	74
	"slots"	72
	"experiences"	63
	"jobs_keywords"	54
	"auth_identities"	45
	"questions"	37
	"educations"	37

4. Yes URL redirection is taking place.

```
>>> redirection=DF_router.loc[DF_router["status"]=="302"]
>>> redirection
```

	at	method	...	code	desc
443	info	GET	...	NaN	NaN
850	info	GET	...	NaN	NaN
869	info	POST	...	NaN	NaN
961	info	GET	...	NaN	NaN
1049	info	GET	...	NaN	NaN
1281	info	GET	...	NaN	NaN
1542	info	GET	...	NaN	NaN
1631	info	GET	...	NaN	NaN
1674	info	GET	...	NaN	NaN
1900	info	GET	...	NaN	NaN
1954	info	GET	...	NaN	NaN
2084	info	POST	...	NaN	NaN

5. Yes, for example the error with code=H18, which according to heroku's documentation, it stands for "Server Request Interrupted" and signifies that the socket connected, some data was sent as part of a response by the app, but then the socket was destroyed without completing the response.

```
>>> server_error=DF_router.loc[DF_router["status"].str.contains("50")]
>>> server_error[["status","code"]]
   status code
1681    503  H18
4206    500  NaN
4209    500  NaN
>>>
```

## **TASK 2**

1.

**SQL\_query=** *SELECT c.last\_name,c.first\_name, c.store\_id, count(r.rental\_id)*  
*FROM customer AS c INNER JOIN rental AS r ON c.customer\_id=r.customer\_id*  
*GROUP BY c.customer\_id HAVING c.store\_id=2 ORDER BY count(rental\_id)*  
*DESC LIMIT 1;*

**Answer=**Seal Karl has the most rentals at store 2.



The screenshot shows a PostgreSQL query editor interface. The top bar indicates the connection is 'dvdrental/postgres@PostgreSQL'. Below the bar are tabs for 'Query Editor' and 'Query History'. The 'Query Editor' tab is active, displaying a SQL query. Below the query editor are tabs for 'Data Output', 'Explain', 'Messages', and 'Notifications'. The 'Data Output' tab is active, showing the results of the query in a table format.

Query:

```
1 SELECT c.last_name,c.first_name, c.store_id, count(r.rental_id)
2 FROM customer AS c
3 INNER JOIN rental AS r
4 ON c.customer_id=r.customer_id
5 GROUP BY c.customer_id
6 HAVING c.store_id=2
7 ORDER BY count(rental_id) DESC LIMIT 1;
```

Results:

	last_name character varying (45)	first_name character varying (45)	store_id smallint	count bigint
1	Seal	Karl	2	45

2.

**SQL\_query=** *select f.title, i.inventory\_id, r.rental\_date, r.return\_date FROM film AS f INNER JOIN inventory AS i ON f.film\_id=i.film\_id LEFT JOIN rental AS r ON i.inventory\_id=r.inventory\_id WHERE f.title='Image Princess' AND (r.rental\_date>'2005-07-29' OR '2005-07-29' >r.return\_date) ORDER BY (r.rental\_date,r.return\_date);*

**Answer=**Yes he would be able to rent the copy of “Image Princess” with inventory id=2092, under the assumption that both stores have access to all inventory\_ids(copys).

dvdrental/postgres@PostgreSQL

Query Editor

Query History

```
1 select f.title, i.inventory_id, r.rental_date, r.return_date
2 FROM film AS f
3 INNER JOIN inventory As i
4 ON f.film_id=i.film_id
5 LEFT JOIN rental AS r
6 ON i.inventory_id=r.inventory_id
7 WHERE f.title='Image Princess' AND (r.rental_date>'2005-07-29' OR '2005-07-29' > r.return_date)
8 ORDER BY (r.rental_date,r.return_date);
9
```

Data Output

Explain

Messages

Notifications

	<div>title</div> <div>character varying (255)</div>	<div>inventory_id</div> <div>integer</div>	<div>rental_date</div> <div>timestamp without time zone</div>	<div>return_date</div> <div>timestamp without time zone</div>	
1	Image Princess	2091	2005-07-06 21:15:38	2005-07-15 00:01:38	
2	Image Princess	2090	2005-07-06 22:08:53	2005-07-07 23:21:53	
3	Image Princess	2089	2005-07-08 00:22:06	2005-07-16 20:16:06	
4	Image Princess	2092	2005-07-10 15:16:30	2005-07-11 14:02:30	
5	Image Princess	2092	2005-08-01 21:11:54	2005-08-09 21:00:54	
6	Image Princess	2091	2005-08-18 05:16:28	2005-08-22 10:32:28	
7	Image Princess	2089	2005-08-21 00:27:46	2005-08-22 22:53:46	
8	Image Princess	2090	2005-08-21 08:00:40	2005-08-27 06:52:40	
9	Image Princess	2092	2005-08-22 11:34:43	2005-08-23 16:52:43	
10	Image Princess	2093	2005-08-23 17:57:28	2005-08-29 20:03:28	

3.

**SQL\_query=** *SELECT EXTRACT(YEAR FROM rental\_date) as yy,EXTRACT(MONTH FROM rental\_date) as mm, COUNT(DISTINCT customer\_id) from rental GROUP BY yy,mm;*

**Answer=**The counts are presented in the picture below.

dvdrental/postgres@PostgreSQL				
Query Editor Query History				
<pre> 1 SELECT EXTRACT(YEAR FROM rental_date) as yy,EXTRACT(MONTH FROM rental_date) as mm, 2 COUNT(DISTINCT customer_id) from rental GROUP BY yy,mm; </pre>				
Data Output Explain Messages Notifications				
	yy double precision	mm double precision	count bigint	
1	2005		5	520
2	2005		6	590
3	2005		7	599
4	2005		8	599
5	2006		2	158

4.

**SQL\_query=** *SELECT COUNT(r.rental\_id), c.name FROM category as c INNER JOIN film\_category AS fc ON c.category\_id=fc.category\_id INNER JOIN film As f ON fc.film\_id=f.film\_id INNER JOIN inventory AS i ON f.film\_id=i.film\_id INNER JOIN rental AS r ON i.inventory\_id=r.inventory\_id GROUP BY c.name ORDER BY COUNT(r.rental\_id) DESC;*

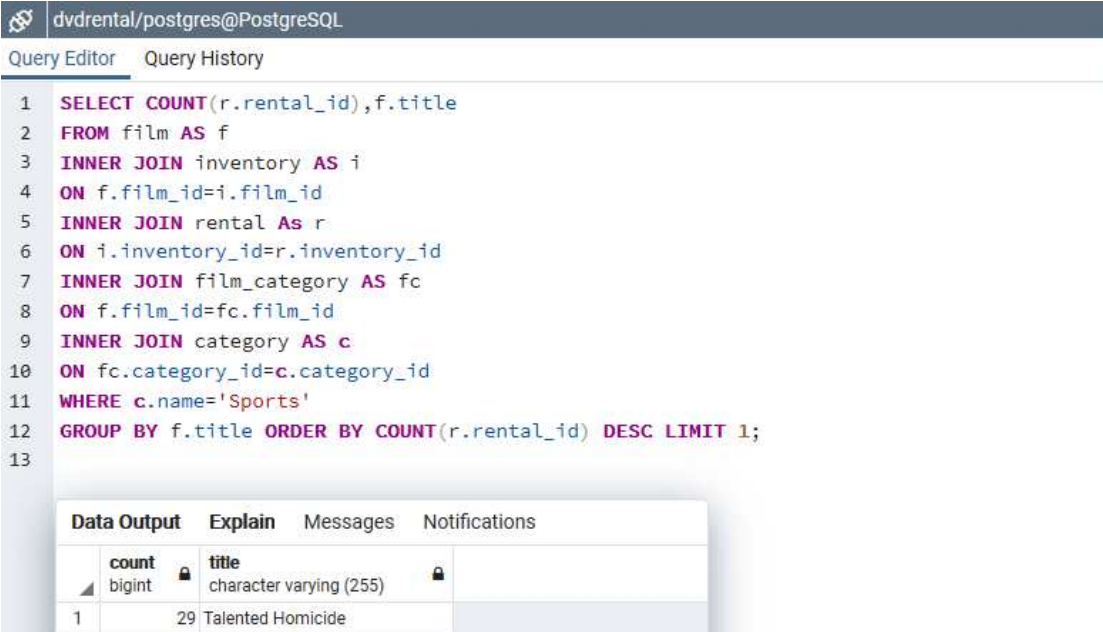
**Answer=**Sports is the most popular category among the store's customers.

dvdrental/postgres@PostgreSQL				
Query Editor Query History				
<pre> 1 SELECT COUNT(r.rental_id), c.name 2 FROM category as c 3 INNER JOIN film_category AS fc 4 ON c.category_id=fc.category_id 5 INNER JOIN film As f ON fc.film_id=f.film_id 6 INNER JOIN inventory AS i ON f.film_id=i.film_id 7 INNER JOIN rental AS r ON i.inventory_id=r.inventory_id 8 GROUP BY c.name ORDER BY COUNT(r.rental_id) DESC; </pre>				
Data Output Explain Messages Notifications				
	count bigint	name character varying (25)		
1	1179	Sports		
2	1166	Animation		
3	1112	Action		
4	1101	Sci-Fi		
5	1096	Family		
6	1060	Drama		
7	1050	Documentary		
8	1033	Foreign		
9	969	Games		
10	945	Children		
11	941	Comedy		
12	940	New		
13	939	Classics		
14	846	Horror		
15	837	Travel		
16	830	Music		

5.

**SQL\_query=** *SELECT COUNT(r.rental\_id),f.title FROM film AS f INNER JOIN inventory AS I ON f.film\_id=i.film\_id INNER JOIN rental As r ON i.inventory\_id=r.inventory\_id INNER JOIN film\_category AS fc ON f.film\_id=fc.film\_id INNER JOIN category AS c ON fc.category\_id=c.category\_id WHERE c.name='Sports' GROUP BY f.title ORDER BY COUNT(r.rental\_id) DESC LIMIT 1;*

**Answer=**Talented Homicide.



The screenshot shows a PostgreSQL query editor interface. At the top, the connection is labeled 'dvdrental/postgres@PostgreSQL'. Below this, there are tabs for 'Query Editor' and 'Query History'. The 'Query Editor' tab is active, displaying the following SQL query:

```
1 SELECT COUNT(r.rental_id),f.title
2 FROM film AS f
3 INNER JOIN inventory AS i
4 ON f.film_id=i.film_id
5 INNER JOIN rental As r
6 ON i.inventory_id=r.inventory_id
7 INNER JOIN film_category AS fc
8 ON f.film_id=fc.film_id
9 INNER JOIN category AS c
10 ON fc.category_id=c.category_id
11 WHERE c.name='Sports'
12 GROUP BY f.title ORDER BY COUNT(r.rental_id) DESC LIMIT 1;
13
```

Below the query editor, there is a 'Data Output' tab which is active, showing the results of the query. The output is a table with two columns: 'count' (bigint) and 'title' (character varying (255)). The first row shows a count of 29 for the title 'Talented Homicide'.

count	title
29	Talented Homicide

6. Examples of other insights that we can obtain from the data are the following:

- Find out which customers are the oldest and more active in order to reward them with a discount.
- Find out which customers exceed the movies' rental duration , in order to impose a penalty.
- Find out for each individual customer what their favorite category is and recommend them the highest rated movies from that category that they haven't already seen.

These actions could increase the business' income and ensure the existence of a loyal customer base.

### **TASK 3**

Web Application using the Flask framework for accessing TMDB API and storing information into MySQL Server.

The project contains 3 different files:

app.py : script for creating and launching the application.

index.html: html page to display the list of movies currently in theatres in Greece

db.yaml: configuration file for connection with db.

#### DDL Commands

CREATE DATABASE movies;

CREATE TABLE now\_playing(movie\_id varchar(30),original\_title varchar(30),title varchar(30),overview text(300));

ALTER TABLE now\_playing MODIFY overview text(1000);

```
mysql> SHOW COLUMNS FROM movies.now_playing;
+-----+-----+-----+-----+-----+-----+
| Field          | Type          | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| movie_id       | varchar(30)   | YES  |     | NULL    |       |
| original_title | varchar(30)   | YES  |     | NULL    |       |
| title          | varchar(30)   | YES  |     | NULL    |       |
| overview       | text          | YES  |     | NULL    |       |
+-----+-----+-----+-----+-----+-----+
4 rows in set (1.21 sec)

mysql>
```



Hi!

ID	Original Title	Title	Overview
451184	Wasp Network	Wasp Network	Havana, Cuba, 1990. René González, an airplane pilot, unexpectedly flees the country, leaving behind his wife Olga and his daughter Irma, and begins a new life in Miami, where he becomes a member of an anti-Castro organization.
522098	Babyteeth	Babyteeth	A terminally ill teen upsets her parents when she falls in love with a small-time drug dealer.
656279	Pari	Pari	Babak, an Iranian student in Greece, doesn't show up to welcome his visiting parents at the Athens airport. Pari and her older husband, both devout Muslims abroad for the first time, are ill-prepared to search for their son in an intimidating and alien environment. All their attempts to find a clue that might lead them to him prove to be in vain and they soon reach a dead end. But Pari can't give up looking for him, even when returning to Iran seems like her only choice. Following the steps of her rebellious son in the darkest corners of the city, she will exhaust her inner strength to achieve more than a mother's search for her missing son.